

Session 9:

Advance Hive

Assignment 1

- By Prachi Mohite

DATE SET DESCRIPTION

The data set consists of the following fields.

Athlete: This field consists of the athlete name

Age: This field consists of athlete ages

Country: This fields consists of the country names which participated in Olympics

Year: This field consists of the year

Closing Date: This field consists of the closing date of ceremony

Sport: Consists of the sports name

Gold Medals: No. of Gold medals

Silver Medals: No. of Silver medals

Bronze Medals: No. of Bronze medals

Total Medals: Consists of total no. of medals

Solution Approach

Here we will be using the database named as 'custom' created in HIVE assignment 8

```
hive> show databases;  
OK  
acadgildb  
custom  
default  
Time taken: 0.167 seconds, Fetched: 3 row(s)  
hive> use custom;  
OK  
Time taken: 0.057 seconds
```

Creating a table

```
CREATE TABLE OLYMPICDATA (  
  Athlete STRING,  
  Age TINYINT,  
  Country STRING,  
  Year INT,  
  Closing_Date STRING,  
  Sports STRING,  
  Gold_Medals TINYINT,  
  Silver_Medals TINYINT,  
  Bronze_Medals TINYINT,  
  Total_Medals SMALLINT) row format delimited fields terminated by '\t'
```

STORED As TEXTFile;

```
hive> CREATE TABLE OLYMPICDATA (  
  > Athlete STRING,  
  > Age TINYINT,  
  > Country STRING,  
  > Year INT,  
  > Closing_Date STRING,  
  > Sports STRING,  
  > Gold_Medals TINYINT,  
  > Silver_Medals TINYINT,  
  > Bronze_Medals TINYINT,  
  > Total_Medals SMALLINT) row format delimited fields terminated by '\t'  
  > STORED As TEXTFile;  
OK  
Time taken: 0.489 seconds  
hive>
```

```
hive> show tables;  
OK  
csv_dump  
olympicdata  
temperature_data  
temperature_data1  
temperature_data_vw  
Time taken: 0.124 seconds, Fetched: 5 row(s)  
hive>
```

Loading the data into newly created table named as olympicdata

LOAD DATA LOCAL INPATH

'/home/acadgild/Desktop/Prachi/HIVE_DATA/olympix_data.csv' INTO TABLE olympicdata;

```
hive> LOAD DATA LOCAL INPATH '/home/acadgild/Desktop/Prachi/HIVE_DATA/olympix_data.csv'  
  > INTO TABLE olympicdata;  
Loading data to table custom.olympicdata  
OK  
Time taken: 2.073 seconds  
hive>
```

Verify if data is loaded properly or not

- With simple select command

```
hive> select * from olympicdata;
OK
Michael Phelps 23 United States 2008 08-24-08 Swimming 8 0 0 8
Michael Phelps 19 United States 2004 08-29-04 Swimming 6 0 2 8
Michael Phelps 27 United States 2012 08-12-12 Swimming 4 2 0 6
Natalie Coughlin 25 United States 2008 08-24-08 Swimming 1 2 3 6
Aleksy Nemov 24 Russia 2000 10-01-00 Gymnastics 2 1 3 6
Alicia Coutts 24 Australia 2012 08-12-12 Swimming 1 3 1 5
Missy Franklin 17 United States 2012 08-12-12 Swimming 4 0 1 5
Ryan Lochte 27 United States 2012 08-12-12 Swimming 2 2 1 5
Allison Schmitt 22 United States 2012 08-12-12 Swimming 3 1 1 5
Natalie Coughlin 21 United States 2004 08-29-04 Swimming 2 2 1 5
Ian Thorpe 17 Australia 2000 10-01-00 Swimming 3 2 0 5
Dara Torres 33 United States 2000 10-01-00 Swimming 2 0 3 5
Cindy Klassen 26 Canada 2006 02-26-06 Speed Skating 1 2 5
Nastia Liukin 18 United States 2008 08-24-08 Gymnastics 1 3 1 5
Marit Bjørgen 29 Norway 2010 02-28-10 Cross Country Skiing 3 1 1 5
Sun Yang 20 China 2012 08-12-12 Swimming 2 1 4
Kirsty Coventry 24 Zimbabwe 2008 08-24-08 Swimming 1 3 0 4
Libby Lenton-Trickett 23 Australia 2008 08-24-08 Swimming 2 1 1 4
Ryan Lochte 24 United States 2008 08-24-08 Swimming 2 0 2 4
Inge de Bruijn 30 Netherlands 2004 08-29-04 Swimming 1 1 2 4
Petria Thomas 28 Australia 2004 08-29-04 Swimming 3 1 0 4
Ian Thorpe 21 Australia 2004 08-29-04 Swimming 2 1 1 4
Inge de Bruijn 27 Netherlands 2000 10-01-00 Swimming 3 1 0 4
Gary Hall Jr. 25 United States 2000 10-01-00 Swimming 2 1 1 4
Michael Klim 23 Australia 2000 10-01-00 Swimming 2 2 0 4
Susie O'Neill 27 Australia 2000 10-01-00 Swimming 1 3 0 4
Jenny Thompson 27 United States 2000 10-01-00 Swimming 3 0 1 4
Pieter van den Hoogenband 22 Netherlands 2000 10-01-00 Swimming 2 0 2 4
An Hyeon-Su 20 South Korea 2006 02-26-06 Short-Track Speed Skating 3 0 1 4
Aliya Mustafina 17 Russia 2012 08-12-12 Gymnastics 1 1 2 4
Shawn Johnson 16 United States 2008 08-24-08 Gymnastics 1 3 0 4
Dmitry Sautin 26 Russia 2000 10-01-00 Diving 1 2 4
Leontien Zijlaard-van Moorsel 30 Netherlands 2000 10-01-00 Cycling 3 1 0 4
Petter Northug Jr. 24 Norway 2010 02-28-10 Cross Country Skiing 2 1 1 4
Ole Einar Bjørndalen 28 Norway 2002 02-24-02 Biathlon 4 0 0 4
Janica Kostelic 20 Croatia 2002 02-24-02 Alpine Skiing 3 1 0 4
Nathan Adrian 23 United States 2012 08-12-12 Swimming 2 1 0 3
Yannick Agnel 20 France 2012 08-12-12 Swimming 2 1 0 3
Brittany Elmslie 18 Australia 2012 08-12-12 Swimming 1 2 0 3
Matt Grevers 27 United States 2012 08-12-12 Swimming 2 1 0 3
Ryosuke Irie 22 Japan 2012 08-12-12 Swimming 0 2 1 3
Cullen Jones 28 United States 2012 08-12-12 Swimming 1 2 0 3
Ranomi Kromowidjojo 21 Netherlands 2012 08-12-12 Swimming 2 1 0 3
```

From HDFS

```
[acagild@localhost ~]$ hadoop fs -ls /user/hive/warehouse/custom.db
18/05/10 04:12:28 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
drwxr-xr-x - acagild supergroup 0 2018-05-10 04:10 /user/hive/warehouse/custom.db/olympicdata
drwxr-xr-x - acagild supergroup 0 2018-05-07 17:09 /user/hive/warehouse/custom.db/temperature_data
drwxr-xr-x - acagild supergroup 0 2018-05-07 17:14 /user/hive/warehouse/custom.db/temperature_data1
[acagild@localhost ~]$ hadoop fs -ls /user/hive/warehouse/custom.db/olympicdata
18/05/10 04:12:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rwxr-xr-x 1 acagild supergroup 518669 2018-05-10 04:10 /user/hive/warehouse/custom.db/olympicdata/olympix_data.csv
```

```
[acadgild@localhost ~]$ hadoop fs -cat /user/hive/warehouse/custom.db/olympicdata/olympix_data.csv
18/05/10 04:13:47 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Michael Phelps 23 United States 2008 08-24-08 Swimming 8 0 0 8
Michael Phelps 19 United States 2004 08-29-04 Swimming 6 0 2 8
Michael Phelps 27 United States 2012 08-12-12 Swimming 4 2 0 6
Natalie Coughlin 25 United States 2008 08-24-08 Swimming 1 2 3 6
Aleksy Nemov 24 Russia 2000 10-01-00 Gymnastics 2 1 3 6
Alicia Coutts 24 Australia 2012 08-12-12 Swimming 1 3 1 5
Missy Franklin 17 United States 2012 08-12-12 Swimming 4 0 1 5
Ryan Lochte 27 United States 2012 08-12-12 Swimming 2 2 1 5
Allison Schmitt 22 United States 2012 08-12-12 Swimming 3 1 1 5
Natalie Coughlin 21 United States 2004 08-29-04 Swimming 2 2 1 5
Ian Thorpe 17 Australia 2000 10-01-00 Swimming 3 2 0 5
Dara Torres 33 United States 2000 10-01-00 Swimming 2 0 3 5
Cindy Klassen 26 Canada 2006 02-26-06 Speed Skating 1 2 2 5
Nastia Liukin 18 United States 2008 08-24-08 Gymnastics 1 3 1 5
Marit Bjørgen 29 Norway 2010 02-28-10 Cross Country Skiing 3 1 1 5
Sun Yang 20 China 2012 08-12-12 Swimming 2 1 1 4
Kirsty Coventry 24 Zimbabwe 2008 08-24-08 Swimming 1 3 0 4
Libby Lenton-Trickett 23 Australia 2008 08-24-08 Swimming 2 1 1 4
Ryan Lochte 24 United States 2008 08-24-08 Swimming 2 0 2 4
Inge de Bruijn 30 Netherlands 2004 08-29-04 Swimming 1 1 2 4
Petria Thomas 28 Australia 2004 08-29-04 Swimming 3 1 0 4
Ian Thorpe 21 Australia 2004 08-29-04 Swimming 2 1 1 4
Inge de Bruijn 27 Netherlands 2000 10-01-00 Swimming 3 1 0 4
Gary Hall Jr. 25 United States 2000 10-01-00 Swimming 2 1 1 4
Michael Klim 23 Australia 2000 10-01-00 Swimming 2 2 0 4
Susie O'Neill 27 Australia 2000 10-01-00 Swimming 1 3 0 4
Jenny Thompson 27 United States 2000 10-01-00 Swimming 3 0 1 4
Pieter van den Hoogenband 22 Netherlands 2000 10-01-00 Swimming 2 0 2 4
An Hyeon-Su 20 South Korea 2006 02-26-06 Short-Track Speed Skating 3 0 1 4
Aliya Mustafina 17 Russia 2012 08-12-12 Gymnastics 1 1 2 4
Shawn Johnson 16 United States 2008 08-24-08 Gymnastics 1 3 0 4
Dmitry Sautin 26 Russia 2000 10-01-00 Diving 1 1 2 4
Leontien Zijlaard-van Moorsel 30 Netherlands 2000 10-01-00 Cycling 3 1 0 4
Petter Northug Jr. 24 Norway 2010 02-28-10 Cross Country Skiing 2 1 1 4
Ole Einar Bjørndalen 28 Norway 2002 02-24-02 Biathlon 4 0 0 4
Janica Kostelic 20 Croatia 2002 02-24-02 Alpine Skiing 3 1 0 4
Nathan Adrian 23 United States 2012 08-12-12 Swimming 2 1 0 3
Yannick Agnel 20 France 2012 08-12-12 Swimming 2 1 0 3
Brittany Elmslie 18 Australia 2012 08-12-12 Swimming 1 2 0 3
Matt Grevers 27 United States 2012 08-12-12 Swimming 2 1 0 3
Ryosuke Irie 22 Japan 2012 08-12-12 Swimming 0 2 1 3
Cullen Jones 28 United States 2012 08-12-12 Swimming 1 2 0 3
Ranomi Kromowidjojo 21 Netherlands 2012 08-12-12 Swimming 2 1 0 3
```

Task 1

1.1 Write a Hive program to find the number of medals won by each country in swimming.

Solution Approach

- Group on Country and get sum of medals
- Where Clause on Sports column to get swimming medals

```
hive>
> select country, sum(total_medals)
> from olympicdata
> where Sports = "Swimming"
> group by Country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine
ive 1.X releases.
Query ID = acadgild_20180510042002_1cda218b-5f78-4361-bf77-305e3e762a61
```

Output

```

MapReduce Total cumulative CPU time: 9 seconds 270 msec
Ended Job = job_1525907687921_0001
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.27 sec HDFS Read: 529303 HDFS Write: 881 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 270 msec
OK
Argentina      1
Australia      163
Austria 3
Belarus 2
Brazil 8
Canada 5
China 35
Costa Rica 2
Croatia 1
Denmark 1
France 39
Germany 32
Great Britain 11
Hungary 9
Italy 16
Japan 43
Lithuania 1
Netherlands 46
Norway 2
Poland 3
Romania 6
Russia 20
Serbia 1
Slovakia 2
Slovenia 1
South Africa 11
South Korea 4
Spain 3
Sweden 9
Trinidad and Tobago 1
Tunisia 3
Ukraine 7
United States 267
Zimbabwe 7
Time taken: 123.082 seconds, Fetched: 34 row(s)
hive> █

```

1.2. Write a Hive program to find the number of medals that India won year wise.

Solution Approach

- Group on Year
- Where Clause on Country column to get India's medals

```

hive> select Year,sum(Total_Medals)
> from olympicdata
> where Country= "India"
> group by Year Order by Year;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future
ive 1.X releases.
Query ID = acadgild_20180510045858_e206e94d-0844-4060-bf43-6fc3cc42ac9a
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
To control the number of reducers, use:
  set hive.exec.parallel=<boolean>

```

Output

```
2018-05-10 05:01:17,373 Stage-2 map = 0%, reduce = 0%
2018-05-10 05:01:31,915 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.06 sec
2018-05-10 05:01:47,568 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 5.6 sec
MapReduce Total cumulative CPU time: 5 seconds 600 msec
Ended Job = job_1525907687921_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.68 sec HDFS Read: 528470 HDFS Write: 180 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 5.6 sec HDFS Read: 5695 HDFS Write: 163 SUCCESS
Total MapReduce CPU Time Spent: 15 seconds 280 msec
OK
2000 1
2004 1
2008 3
2012 6
Time taken: 170.971 seconds, Fetched: 4 row(s)
hive>
```

1.3. Write a Hive Program to find the total number of medals each country won.

Solution Approach

- Group on Country and sum of total medals

```
hive> select country, sum(Total_Medals)
> from olympicdata
> group by Country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future
ive 1.X releases.
Query ID = acadgild_20180510050257_a599d020-1792-464c-9d30-6d2b9e67b397
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>
To order to limit the maximum number of reducers:
```

Output

```

Total MapReduce CPU Time Spent: 7 seconds 200 msec
OK
Afghanistan      2
Algeria          8
Argentina        141
Armenia          10
Australia        609
Austria          91
Azerbaijan       25
Bahamas          24
Bahrain          1
Barbados         1
Belarus          97
Belgium          18
Botswana         1
Brazil           221
Bulgaria         41
Cameroon         20
Canada           370
Chile            22
China            530
Chinese Taipei   20
Colombia         13
Costa Rica       2
Croatia          81
Cuba             188
Cyprus           1
Czech Republic   81
Denmark          89
Dominican Republic 5
Ecuador          1
Egypt            8
Eritrea          1
Estonia          18
Ethiopia         29
Finland          118
France           318
Gabon            1
Georgia          23
Germany          629
Great Britain    322
Greece           59

```

1.4. Write a Hive program to find the number of gold medals each country won.

Solution Approach

- Group on Country and get sum of gold medals

```

hive>
> select country, sum(Gold_Medals)
> from olympicdata
> group by Country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in t
ive 1.X releases.
Query ID = acadgild_20180510064933_9ee2906b-8f16-401c-b448-728d2ed6a3d6
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:

```


Output

```
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.42 sec HDFS Read: 528469 HDFS Write: 2703 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 420 msec
OK
Afghanistan 0
Algeria 2
Argentina 49
Armenia 0
Australia 163
Austria 36
Azerbaijan 6
Bahamas 11
Bahrain 0
Barbados 0
Belarus 17
Belgium 2
Botswana 0
Brazil 46
Bulgaria 8
Cameroon 20
Canada 168
Chile 3
China 234
Chinese Taipei 2
Colombia 2
Costa Rica 0
Croatia 35
Cuba 57
Cyprus 0
Czech Republic 14
Denmark 46
Dominican Republic 3
Ecuador 0
Egypt 1
Eritrea 0
Estonia 6
Ethiopia 13
Finland 11
France 108
Gabon 0
Georgia 6
Germany 223
Great Britain 124
Greece 12
Grenada 1
Guatemala 0
```

Task 2

Write a hive UDF that implements functionality of string concat_ws(string SEP, array<string>). This UDF will accept two arguments, one string and one array of string. It will return a single string where all the elements of the array are separated by the SEP.

UDFs

User Defined Functions, also known as UDF, allow you to create custom functions to process records or groups of records.

Hive has

- UDF
 - A UDF processes one or several columns of one row and outputs one value.
 - UDAF - User-Defined Aggregation Functions
 - Aggregate functions perform a calculation on a set of values and return a single value. An aggregate function is more difficult to write than a regular UDF. Values are aggregated in chunks (potentially across many tasks), so the implementation has to be capable of combining partial aggregations into a final result.
 - Simple UDF
 - Where all complex datatypes cannot be handled
 - Reduced performance due to use of reflection: each call of the evaluate method is reflective. Furthermore, all arguments are evaluated and parsed.
 - Generic UDF
 - A generic UDF is written by extending the GenericUDF class.
 - All complex parameters are supported (even nested ones like array<array>)
 - Variable number of arguments are supported

Solution Approach

Here we will create Generic UDF as Generic UDF allows to accept Complex Datatypes like Array of string

- To create a GenericUDF, need to inherit the abstract class '**GenericUDF**'
- Need to overwrite below methods
- **ObjectInspector initialize(ObjectInspector[] arguments)**
 - To verify what type of data types / arguments are accepted in UDF
 - Set up and return an ObjectInspector for the type of the output of the UDF
 - Store in global variables the ObjectInspectors for the elements of the input
 - Set up the storage variable for the output
- **Object evaluate(DeferredObject[] arguments)**
 - Actual functionality of generic UDF which should get executed once UDF is called through hive shell / script
- **String getDisplayString(String[] children)**
 - returns the string that will be returned when explain is used

```

1 import org.apache.hadoop.hive.q1.exec.TextRecordWriter;
2 import org.apache.hadoop.hive.q1.exec.UDFArgumentException;
3 import org.apache.hadoop.hive.q1.exec.UDFArgumentLengthException;
4 import org.apache.hadoop.hive.q1.exec.UDFArgumentTypeException;
5 import org.apache.hadoop.hive.q1.metadata.HiveException;
6 import org.apache.hadoop.hive.q1.udf.generic.GenericUDF;
7 import org.apache.hadoop.hive.serde.serdeConstants;
8 import org.apache.hadoop.hive.serde2.objectinspector.ListObjectInspector;
9 import org.apache.hadoop.hive.serde2.objectinspector.ObjectInspector;
10 import org.apache.hadoop.hive.serde2.objectinspector.ObjectInspector.Category;
11 import org.apache.hadoop.hive.serde2.objectinspector.PrimitiveObjectInspector;
12 import org.apache.hadoop.hive.serde2.objectinspector.PrimitiveObjectInspector.PrimitiveCategory;
13 import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorFactory;
14 import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorUtils;
15 import org.apache.hadoop.hive.serde2.objectinspector.primitive.PrimitiveObjectInspectorUtils.PrimitiveGrouping;
16 import org.apache.hadoop.io.*;
17
18 public class Concat_WS extends GenericUDF {
19     private transient ObjectInspector[] argumentOIs;
20     @Override
21     public ObjectInspector initialize(ObjectInspector[] arguments) throws
22     UDFArgumentException {
23
24         if (arguments.length < 2) {
25             throw new UDFArgumentLengthException(
26                 "The function CONCAT_WS(separator,[string | array(string)]+) "
27                 + "needs at least two arguments.");
28         }
29         // check if argument is a string or an array of strings
30         for (int i = 0; i < arguments.length; i++) {
31             switch(arguments[i].getCategory()) {
32                 case LIST:
33                     if (isStringOrVoidType(
34                         ((ListObjectInspector) arguments[i]).getListElementObjectInspector())) {
35                         break;
36                     }
37                     case PRIMITIVE:
38
39                         if (isStringOrVoidType(arguments[i])) {
40                             break;
41                         }
42                     default:
43                         throw new UDFArgumentTypeException(i, "Argument " + (i + 1)
44                             + " of function CONCAT_WS must be \"" + serdeConstants.STRING_TYPE_NAME
45                             + " or " + serdeConstants.LIST_TYPE_NAME + "<" +
46                             serdeConstants.STRING_TYPE_NAME
47                             + ">\", but \"" + arguments[i].getTypeName() + "\" was found.");
48             }
49         }
50         argumentOIs = arguments;
51         return PrimitiveObjectInspectorFactory.writableStringObjectInspector;
52     }
53
54     protected boolean isStringOrVoidType(ObjectInspector oi) {
55         if (oi.getCategory() == Category.PRIMITIVE) {
56             if (PrimitiveGrouping.STRING_GROUP
57                 == PrimitiveObjectInspectorUtils.getPrimitiveGrouping(
58                     ((PrimitiveObjectInspector) oi).getPrimitiveCategory())
59                 || ((PrimitiveObjectInspector) oi).getPrimitiveCategory() == PrimitiveCategory.VOID)
60             {
61                 return true;
62             }
63         }
64         return false;
65     }
66
67     private final Text resultText = new Text();
68
69     @Override
70     public Object evaluate(DeferredObject[] arguments) throws HiveException {
71         if (arguments[0].get() == null) {
72             return null;
73         }
74         String separator = PrimitiveObjectInspectorUtils.getString(

```

```

arguments[0].get(), (PrimitiveObjectInspector)argumentOIs[0]);
StringBuilder sb = new StringBuilder();
boolean first = true;
for (int i = 1; i < arguments.length; i++) {
    if (arguments[i].get() != null) {
        if (first) {
            first = false;
        } else {
            sb.append(separator);
        }
        if (argumentOIs[i].getCategory().equals(Category.LIST)) {
            Object strArray = arguments[i].get();
            ListObjectInspector strArrayOI = (ListObjectInspector) argumentOIs[i];
            boolean strArrayFirst = true;
            for (int j = 0; j < strArrayOI.getListLength(strArray); j++) {
                if (strArrayFirst) {
                    strArrayFirst = false;
                } else {
                    sb.append(separator);
                }
                sb.append(strArrayOI.getListElement(strArray, j));
            }
        } else {
            sb.append(PrimitiveObjectInspectorUtils.getString(
                arguments[i].get(), (PrimitiveObjectInspector)argumentOIs[i]));
        }
    }
}
resultText.set(sb.toString());
return resultText;
}

@Override
public String getDisplayString(String[] children) {
    assert (children.length >= 2);
    return getStandardDisplayString("concat_ws", children);
}

```

```

hive> add jar /home/acadgild/Desktop/Prachi/HIVE_UDF/HIVE_UDF.jar;
Added /home/acadgild/Desktop/Prachi/HIVE_UDF/HIVE_UDF.jar to class path
Added resources: [/home/acadgild/Desktop/Prachi/HIVE_UDF/HIVE_UDF.jar]
hive> create temporary function Concat_WS as 'Concat_WS';
OK
Time taken: 1.101 seconds
hive> select Concat_WS('-', 'SAYALI', 'Prachi', 'Aboli');
OK
SAYALI-Prachi-Aboli
Time taken: 2.844 seconds, Fetched: 1 row(s)

```

Executing UDF on Table which has one column as array of strings

```

hive> create table temp ( date1 string, city string, MyTemp array<string> ) row format delimited fields terminated by '\t' collection items terminated by ',';
OK
Time taken: 0.33 seconds
hive> load data local inpath '/home/acadgild/mytempdata.txt' into table temp;
Loading data to table custom.temp
OK
Time taken: 1.796 seconds
hive> SELECT * FROM TEMP;
OK
1/2/2017      A      ["1","2","3","4","5","6"]
1/2/2017      B      ["1","2","3","4","5","6"]
1/2/2017      C      ["1","2","3","4","5","6"]

```

```

hive> select Concat_WS('-',city,mytemp) from temp;
OK
A-1-2-3-4-5-6
B-1-2-3-4-5-6
C-1-2-3-4-5-6
Time taken: 0.529 seconds, Fetched: 5 row(s)

```

Output from running Concat_WS on table having column as array of string

Task 3

Link: <https://acadgild.com/blog/transactions-in-hive/>

Refer the above given link for transactions in Hive and implement the operations given in the blog using your own sample data set and send us the screenshot.

Transactions are provided at the row-level in Hive 0.14. The different row-level transactions available in Hive 0.14 are as follows:

1. Insert
2. Delete
3. Update

There are numerous limitations with the present transactions available in Hive 0.14. ORC is the file format supported by Hive transaction. It is now essential to have ORC file format for performing transactions in Hive.

The transaction features present in Hive needs to be turned on, as by default they are turned off.

- set hive.support.concurrency = true;
 - o Default value is false
- set hive.enforce.bucketing = true;
 - o Ensures we can create buckets on a table
- set hive.exec.dynamic.partition.mode = nonstrict;
 - o This is to set dynamic partition mode on
- set hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
- set hive.compactor.initiator.on = true;
- set hive.compactor.worker.threads = a positive number on at least one instance of the Thrift metastore service

```
hive> set hive.support.concurrency = true;
hive>
> set hive.enforce.bucketing = true;
hive>
> set hive.exec.dynamic.partition.mode = nonstrict;
hive>
> set hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
hive>
> set hive.compactor.initiator.on = true;
hive>
> set hive.compactor.worker.threads = 10;
hive>
```

Create a Buckted table which will support transaction

```
CREATE TABLE college(clg_id int,clg_name string,clg_loc string) clustered by (clg_id) into 5 buckets
stored as orc TBLPROPERTIES('transactional'='true');
```

```
hive> create database transactions;
OK
Time taken: 0.184 seconds
```

```
hive> use transactions;
OK
Time taken: 0.035 seconds
hive> show tables;
OK
Time taken: 0.085 seconds
hive> CREATE TABLE college(clg_id int,clg_name string,clg_loc string) clustered by (clg_id) into 5 buckets stored as orc TBLPROPERTIES('transactional'='true');
OK
Time taken: 0.392 seconds
hive> show tables;
OK
college
Time taken: 0.075 seconds, Fetched: 1 row(s)
hive> █
```

Inserting values in college table

insert into table college

values(1,'Cummins','Karvenagar'),(2,'PICT','Dhankawadi'),(3,'VIT','Bibewadi'),(4,'COEP','Shivaji Nagar'),(5,'MIT','Kothrud');

The above command is used to insert row wise data into the Hive table. Here, each row is separated by '(',')'

```
hive> insert into table college
> values(1,'Cummins','Karvenagar'),(2,'PICT','Dhankawadi'),(3,'VIT','Bibewadi'),(4,'COEP','Shivaji Nagar'),(5,'MIT','Kothrud');
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180510233204_c2b7bf3-ccd8-44ef-bf96-229779987ef1
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1525968041555_0001, Tracking URL = http://localhost:8088/proxy/application_1525968041555_0001/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1525968041555_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 5
2018-05-10 23:32:45,476 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 5.66 sec
2018-05-10 23:33:07,849 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 6.71 sec
2018-05-10 23:34:03,793 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 6.71 sec
2018-05-10 23:34:05,717 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 11.31 sec
2018-05-10 23:34:39,465 Stage-1 map = 100%, reduce = 78%, Cumulative CPU 33.97 sec
2018-05-10 23:34:42,742 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 36.3 sec
2018-05-10 23:35:17,648 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 39.65 sec
2018-05-10 23:35:31,277 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 45.01 sec
MapReduce Total cumulative CPU time: 45 seconds 10 msec
Ended Job = job_1525968041555_0001
Loading data to table transactions.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 5 Cumulative CPU: 45.01 sec HDFS Read: 27269 HDFS Write: 4094 SUCCESS
Total MapReduce CPU Time Spent: 45 seconds 10 msec
OK
Time taken: 211.296 seconds
hive> select * from college;
OK
5      MIT      Kothrud
1      Cummins  Karvenagar
2      PICT     Dhankawadi
3      VIT      Bibewadi
4      COEP     Shivaji Nagar
Time taken: 0.536 seconds, Fetched: 5 row(s)
hive> █
```

we will re-insert the same data again, it will be appended to the previous data as shown below:

```

In order to set a constant number of reducers:
set mapreduce.job.reduces=<number>
Starting Job = job_1525968041555_0002, Tracking URL = http://localhost:8088/proxy/application_1525968041555_0002/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1525968041555_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 5
2018-05-10 23:37:39,606 Stage-1 map = 0%, reduce = 0%
2018-05-10 23:37:55,202 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.1 sec
2018-05-10 23:38:35,694 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 5.33 sec
2018-05-10 23:38:45,277 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 7.91 sec
2018-05-10 23:38:48,230 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 9.83 sec
2018-05-10 23:38:49,722 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 11.94 sec
2018-05-10 23:39:02,164 Stage-1 map = 100%, reduce = 60%, Cumulative CPU 16.53 sec
2018-05-10 23:39:07,218 Stage-1 map = 100%, reduce = 66%, Cumulative CPU 21.3 sec
2018-05-10 23:39:10,331 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 21.88 sec
2018-05-10 23:39:11,802 Stage-1 map = 100%, reduce = 73%, Cumulative CPU 26.6 sec
2018-05-10 23:39:14,344 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 31.86 sec
2018-05-10 23:39:37,139 Stage-1 map = 100%, reduce = 93%, Cumulative CPU 34.19 sec
2018-05-10 23:39:48,008 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 38.06 sec
MapReduce Total cumulative CPU time: 38 seconds 60 msec
Ended Job = job_1525968041555_0002
Loading data to table transactions.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 5 Cumulative CPU: 38.06 sec HDFS Read: 27054 HDFS Write: 4098 SUCCESS
Total MapReduce CPU Time Spent: 38 seconds 60 msec
OK
Time taken: 153.4 seconds
hive> select * from college;
OK
5      MIT      Kothrud
5      MIT      Kothrud
1      Cummins  Karvenagar
1      Cummins  Karvenagar
2      PICT     Dhankawadi
2      PICT     Dhankawadi
3      VIT      Bibewadi
3      VIT      Bibewadi
4      COEP     Shivaji Nagar
4      COEP     Shivaji Nagar
Time taken: 0.57 seconds, Fetched: 10 row(s)
hive>

```

Updating the Data in Hive Table on bucketed column

Update college set clg_id =20 where clg_id = 5;

```

hive> Update college set clg_id =20 where clg_id = 5;
FAILED: SemanticException [Error 10302]: Updating values of bucketing columns is not supported. Column clg_id.
hive>

```

But we can perform update operation on Non bucketed column

Update college set clg_name ='cummins College Of engineering' where clg_id = 1;

```

hive> Update college set clg_name='cummins College Of engineering' where clg_id = 1;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using H
ive 1.X releases.
Query ID = acadgild_20180510234351_f3fbcf52-25bc-44a2-9c27-5246a65ea18a
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1525968041555_0003, Tracking URL = http://localhost:8088/proxy/application_1525968041555_0003/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1525968041555_0003
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-05-10 23:44:11,258 Stage-1 map = 0%, reduce = 0%
2018-05-10 23:45:11,629 Stage-1 map = 0%, reduce = 0%
2018-05-10 23:47:20,321 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 38.74 sec
2018-05-10 23:48:16,670 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 45.64 sec
2018-05-10 23:48:31,323 Stage-1 map = 100%, reduce = 60%, Cumulative CPU 48.52 sec
2018-05-10 23:48:32,765 Stage-1 map = 100%, reduce = 77%, Cumulative CPU 56.13 sec
2018-05-10 23:48:35,414 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 57.1 sec
2018-05-10 23:48:49,799 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 60.79 sec
MapReduce Total cumulative CPU time: 1 minutes 0 seconds 790 msec
Ended Job = job_1525968041555_0003
Loading data to table transactions.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 60.79 sec HDFS Read: 57413 HDFS Write: 1135 SUCCESS
Total MapReduce CPU Time Spent: 1 minutes 0 seconds 790 msec
OK
Time taken: 384.005 seconds
hive> select * from college where clg_id = 1;
OK
1 cummins College Of engineering Karvenagar
1 cummins College Of engineering Karvenagar
Time taken: 1.682 seconds, Fetched: 2 row(s)
hive>

```

Deleting a Row from Hive Table:

delete from college where clg_id=5;

```

hive> delete from college where clg_id=5;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using
ive 1.X releases.
Query ID = acadgild_20180510235021_47092024-faf4-4533-8c04-f897588bb14a
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1525968041555_0004, Tracking URL = http://localhost:8088/proxy/application_1525968041555_0004/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1525968041555_0004
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2018-05-10 23:50:46,189 Stage-1 map = 0%, reduce = 0%
2018-05-10 23:51:46,709 Stage-1 map = 0%, reduce = 0%
2018-05-10 23:52:53,864 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 16.04 sec
2018-05-10 23:53:23,138 Stage-1 map = 13%, reduce = 0%, Cumulative CPU 16.04 sec
2018-05-10 23:53:28,658 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 33.28 sec
2018-05-10 23:53:36,494 Stage-1 map = 53%, reduce = 0%, Cumulative CPU 40.84 sec
2018-05-10 23:53:49,280 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 56.27 sec
2018-05-10 23:54:49,526 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 56.27 sec
2018-05-10 23:55:10,692 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 58.53 sec
2018-05-10 23:55:13,897 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 63.34 sec
2018-05-10 23:55:26,440 Stage-1 map = 100%, reduce = 60%, Cumulative CPU 64.02 sec
2018-05-10 23:55:28,001 Stage-1 map = 100%, reduce = 73%, Cumulative CPU 72.12 sec
2018-05-10 23:55:29,257 Stage-1 map = 100%, reduce = 80%, Cumulative CPU 74.71 sec
2018-05-10 23:55:45,320 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 78.89 sec
MapReduce Total cumulative CPU time: 1 minutes 18 seconds 890 msec
Ended Job = job_1525968041555_0004
Loading data to table transactions.college
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 78.89 sec HDFS Read: 55261 HDFS Write: 780 SUCCESS
Total MapReduce CPU Time Spent: 1 minutes 18 seconds 890 msec
OK
Time taken: 332.914 seconds
hive> select * from college where clg_id = 5;
OK
Time taken: 2.593 seconds
hive>

```